The Fractal Predictive Cortex: a training-free, alignment-free, self-organising model that unifies prediction, anomaly detection, classification, and phylogenetics for biological sequences

Md Rafiur Rahman[1]

[1] Department of Biochemistry and Biotechnology, North South University, Dhaka, Bangladesh — *corresponding author:* `rafiur2810@gmail.com`

Keywords: alignment-free sequence comparison · Chaos Game Representation · predictive coding · minimum description length · self-organising maps · phylogenetics · anomaly detection · genome classification

Working draft. All numbers and figures in this paper are produced by the accompanying open-source code (`benchmarks/` and `examples/`; see §6), including the cases where the method matches or is outperformed by its baselines.

Abstract

Alignment-free sequence analysis usually treats each task (compression, classification, distance estimation, anomaly detection) as a separate problem solved by a purpose-built tool, and most of these tools depend on a chosen k-mer length k whose best value differs from dataset to dataset. This paper describes the Fractal Predictive Cortex (FPC), a single self-organising model that performs all four tasks without training labels, sequence alignment, pretrained weights, or a fixed k-mer order. FPC streams a sequence through a Chaos Game Representation, so that the recent context becomes a location in the plane, and it incrementally adds predictive "prototypes" at the locations where its own next-symbol prediction is weakest. Its only objective is to minimise cumulative surprise, which equals the sequence's description length under the model; the four capabilities are different read-outs of that one quantity.

On synthetic data with known ground truth, FPC compresses better than gzip and low-order Markov models and stays close to the best hand-tuned fixed-order Markov model without needing an order to be chosen. It classifies short reads down to 150 bp at 100% accuracy, where k-mer + kNN/logistic baselines fall to 94–97%. Across 25 independent random phylogenies its distances correlate with the true patristic distances about as well as the standard tools FFP and Mash and the earlier compression-based distances NRC and NCCD (mean Pearson r = 0.93; paired Wilcoxon $p \geq 0.24$ against all four), and significantly better than a k-mer cosine baseline ($r = 0.78$, $p < 10^{-6}$). For foreign-insert detection it matches a strong k-mer-novelty detector and adds a calibrated per-position surprise signal.

We also test on real data: seven coronaviruses, a 28-genome five-family set, 40 proteins, and 17 mammalian mitochondrial genomes. FPC recovers the expected close relationships without supervision (the SARS pair, the dengue serotypes, the ebolaviruses, HIV/SIV; four of five viral families form clean clades), runs unchanged on the 20-letter protein alphabet, and flags spliced phiX174 contamination (per-window ROC-AUC = 1.0). On the community-standard Robinson–Foulds mitogenome test it matches Mash, FFP and NCCD (normalised RF 0.43) and is slightly behind NRC (0.36), while beating k-mer cosine (0.71). Where the method loses we say so: on the most divergent viral families the composition method

FFP is clearly more accurate (family purity 0.82 vs 0.93), and near-incompressible viral DNA limits compression for every method.

Because the model is itself a geometry, it can be displayed directly as a three-dimensional map of its internal state. A literature review indicates that FPC's individual components (CGR, self-organising maps, prediction- and compression-based comparison, the least-surprised-model classifier, and the surprise/description-length identity) are all established; the contribution is their combination into one model whose memory grows where prediction fails, for which we found no close precedent. The core engine is about 1,300 lines of NumPy and SciPy and runs offline.

---

1. Introduction

A great deal of modern bioinformatics is *alignment-free*: rather than aligning sequences base-by-base (which is expensive and assumes homology), we summarise a sequence by the statistics of its short sub-words (*k*-mers) and compute on those [Vinga & Almeida 2003; Zielezinski et al. 2017]. This is fast and general, but it fragments into task-specific pipelines: one tool compresses, another classifies, another builds trees, another scans for anomalies — and almost all of them hinge on a hand-chosen word length k whose best value is unknown in advance and differs from dataset to dataset.

We instead ask whether a single model can address all of these tasks from one internal signal. FPC combines two established ideas, coupled in a way for which we could not find a precedent (§1.1):

1. **Chaos Game Representation (CGR)** [Jeffrey 1990] maps a sequence to a trajectory of points in the unit square such that *positions sharing a long common suffix land at nearly the same point*. Geometry becomes context.
2. **Predictive coding / surprise minimisation** — the idea, from both data compression (MDL, PPM, context-tree weighting [Rissanen 1978; Cleary & Witten 1984; Willems et al. 1995]) and theoretical neuroscience [Rao & Ballard 1999; Friston 2010], that a good model of a stream is one that predicts the next token well, and that *prediction error (surprise)* is the signal worth acting on.

FPC places a growing, self-organising set of predictive units **in CGR space** and trains them online by surprise: it spends memory (grows a new unit) exactly where its predictions fail. The result is a model that is simultaneously a predictor, a compressor, a novelty detector, a generative classifier, and a distance function — and, because it lives in a 2-D geometry, a picture.

What this paper claims, and what it does not. We do *not* claim state-of-the-art on any single task: a matched fixed-order Markov model beats FPC on Markov-generated compression, and the composition method FFP beats it on deep-divergence real phylogeny — we report those losses plainly. We *do* claim (a) a genuinely unified framework in which four capabilities fall out of one mechanism and one objective; (b) *competitiveness* with standard alignment-free tools — statistically indistinguishable from FFP and Mash on controlled phylogenies (§3.4) and robust on short reads (§3.3); (c) freedom from k-mer-order selection; and (d) a novel, interpretable 3-D visualisation of the model's internal state and a calibrated per-position surprise signal that single-purpose distance tools do not provide.

1.1 Related work and the precise locus of novelty

Two systematic, adversarially-verified literature reviews (one broad, one focused on the hardest residual claims) establish that **every individual ingredient of FPC is prior art**, and pin down exactly what is not:

- **CGR + neural / self-organising models.** CGR/FCGR has been fed to CNNs (*CGRclust*; Alipour, Hill & Kari, 2024) and combined with a static k-mer "dynamical language" weighting for phylogeny (*CGRWDL*; Wang, Yu & Li, 2024). The closest CGR + self-organising-map work is *SOCGR* (Neme et al., 2008) — but it reorders a **fixed** $2^k \times 2^k$ lattice for topological continuity; it does not grow, predict, or use surprise. FPC shares the CGR substrate with these, not the mechanism.
- **The "least-surprised model" classifier is textbook.** Training one predictor per class and labelling a query by `argmax_i P^i(x)` (equivalently, minimum −log P = least surprise) is the winner-take-all variable-order-Markov classifier of Begleiter, El-Yaniv & Yona (2004), demonstrated on SCOP protein folds. We claim no novelty for this rule.
- **Prediction/compression comparison is mature.** Mixture-of-predictor DNA compressors (*GeCo3*, Silva et al. 2020; *Jarvis*, Pratas et al. 2019) and compression distances (NCD/IBD; La Rosa et al. 2013; Ferragina et al. 2007) embody the same prediction-as-modelling philosophy with fixed-topology mixers, not CGR or growing memory.
- **surprise = code length = MDL** is standard information theory (Rissanen; Cover & Thomas), and the FCGR of order k is provably equivalent to the k-mer frequency vector at that order (He, Kari & Millán-Arias, 2025) — so neither the objective nor the representation is itself novel.
- **Growing neural models on sequences already exist** — for clustering, not prediction. A growing self-organising map has been applied to protein-sequence classification (Ahmad, Alahakoon & Chau, 2010) and a neural gas to DNA k-mer barcodes (Fiannaca et al., 2015). Crucially, these grow units on *vector- quantisation* error (or ART vigilance) to cluster fixed feature vectors — **not** on *next-symbol prediction* error, and not for autoregressive prediction.
- **The cross-entropy distance itself is NOT novel.** Defining a genomic distance as the bits to describe x under a predictor trained on y — i.e. the cross-entropy / relative compression of one sequence's learned model on another, used to build phylogenetic trees — is exactly Pratas, Silva & Pinho's **Normalized Relative Compression** (2018) and **NCCD** (Pratas & Pinho, 2014), and the KL-between-per-sequence-Markov-models trees of Bai (2013) and Chang & Wang (2011). We claim **no novelty** for the distance definition, and we benchmark directly against NRC and NCCD (§3.4, §3.9); our only difference is the *substrate* (a growing predictor rather than a finite-context Markov model).

What is left, and is our contribution. After the corrections above, the defensible novelty is narrow and specific: (a) growth driven by **next-symbol prediction error** ("grow-where-surprised") in an *online, autoregressive* predictor of biological sequences — every prior growing model on sequences grows on quantisation error for clustering, not on predictive surprise; (b) an explicit **predictive-coding framing** of sequence modelling, for which the review found no prior genomic/proteomic application; and (c) the **unification** of these with the CGR substrate so that four capabilities and a 3-D interpretable map arise from one surprise signal — no surveyed method combines two or more of these elements. We therefore frame FPC honestly as a novel **unification/architecture with a prediction-error-driven growth rule**, *not* as a new distance, objective, representation, or classifier, and (per the AFproject no-free-lunch finding) make no claim of universal superiority.

---

2. Method

2.1 Notation and the alphabet

A sequence is a string over an alphabet of a canonical symbols (DNA: $a=4$; protein: $a=20$), plus a wildcard for ambiguous characters. Everything below is alphabet-agnostic; only the symbol-to-integer map knows about letters.

2.2 The fractal mind-space (Chaos Game Representation)

Fix a vertex $v(s) \in [0,1]^2$ for each canonical symbol (the four corners of the unit square for DNA; a regular polygon on a circle otherwise). Starting from the centre $p_0 = (½, ½)$, read the sequence and iterate the contraction map

$$p_t = (1 - r) \cdot p_{t-1} + r \cdot v(s_t), \qquad 0 < r < 1.$$

For DNA we use the classical $r = ½$; in general $r = \min(½, \sin(\pi/a))$ keeps the per-symbol sub-copies from overlapping. Because the map is a contraction, the influence of a symbol decays as $(1-r)$ per step, so **p_t is an exponentially-weighted memory of the recent suffix**. Concretely, two positions whose last k symbols agree differ by at most $\sim(1-r)^k$ in position (empirically $0.5^8 \cdot \sqrt{2} \approx 0.0055$ for an 8-symbol shared suffix, versus ~ 0.1 for a different suffix — a >10× separation; see `tests/test_core.py::test_cgr_shared_suffix_are_close`). We call $C_t = p_{t-1}$ the **context point** for position t: the geometric summary of everything the model has read before symbol t.

The same trajectory, binned into a grid, yields the **Frequency-CGR image** — a fixed-length, alignment-free fingerprint of a whole sequence, here generalised to any alphabet.

2.3 The cortex: prototypes that grow where the model is surprised

The model keeps a set of **prototypes** $\{(x_j, n_j, u_j)\}$, where $x_j \in [0,1]^2$ is a location, $n_j \in \mathbb{R}^a$ a tally of which symbol followed when the context landed near x_j, and u_j a (decaying) usage counter.

Prediction. At context point C_t, take the κ nearest prototypes, weight them by a Gaussian of their distance ($w_j \propto \exp(-\|C_t - x_j\|^2 / 2\sigma^2)$, with σ equal to the split resolution below), and form the next-symbol distribution

$$P(s \mid C_t) = (1-\lambda) \cdot \frac{\sum_j w_j n_j[s] + \alpha}{\sum_j w_j \sum_s n_j[s] + a\alpha} + \lambda \cdot P_{bg}(s),$$

i.e. a distance-weighted, Laplace-smoothed (α) mixture of the local prototypes, blended with a small background term (λ) so every probability is strictly positive and surprise is always finite.

Surprise. The model observes the true s_t and pays

$$\text{surprise}_t = -\log_2 P(s_t \mid C_t) \quad \text{[bits]}.$$

Summed over the sequence, this is the code length under the model — the object FPC minimises, i.e. the MDL criterion.

Online learning and the growth decision. After scoring, FPC reinforces the winning prototype (`n_win[s_t] += 1`) and nudges it toward the context (`x_win += η (C_t − x_win)`, SOM-style). Crucially, it *grows a new prototype at* `C_t` when it is in unfamiliar or hard territory:

```
grow  if  nearest-distance > τ          (empty region), or
          surprise_t is abnormally high  (a region it keeps getting wrong),
where τ = r^order.
```

The single knob `order` sets the split resolution `τ = r^order`, which — because CGR distance halves per shared suffix symbol — corresponds to resolving roughly `order` symbols of context. This is the model's *maximum* effective Markov order; it uses that resolution only where surprise justifies it, and coarser resolution elsewhere. A usage-decay + prune step caps total memory. This "grow where you are surprised" rule is the point of contact between the CGR geometry and predictive coding, and is, to our knowledge, the novel element of the method.

2.4 Four read-outs of one signal

- **Predict.** `P(s | C_t)` directly; argmax for next-base calling, sampling for generation, per-symbol log-loss for evaluation.
- **Detect.** The per-position surprise trace, smoothed and converted to a robust z-score (median/MAD) with a per-region Gaussian p-value and a Bonferroni correction; spikes are candidate mutations, inserts, or contamination. A reference brain (trained on a reference genome) flags stretches of a query it cannot explain.
- **Classify.** Train one cortex per class; label a query by *which brain is least surprised* — the minimum cross-entropy / MDL rule, which for equal priors is the Bayes-optimal decision for the learned generative models. Softmax over `−bits` gives calibrated posteriors.
- **Discover.** Between two sequences define the symmetric *excess* cross-entropy `d(A,B) = ½[(CE_A(B) − CE_B(B)) + (CE_B(A) − CE_A(A))]`, a Kullback–Leibler-style divergence of their generative statistics (0 for identical, growing with divergence). Feed it to classical MDS (for a 3-D embedding) or hierarchical clustering (for a tree).

2.5 The 3-D cognitive map

Because the model lives in `[0,1]²`, we visualise it directly: the floor is CGR space, the height/colour is surprise (the **surprise landscape**), and the prototypes float above their territory coloured by the symbol they predict. A foreign insert appears as a bright ridge; a mastered motif as a valley. A second view places whole organisms in 3-D by their cross-entropy distance (MDS) with a minimum-spanning-tree "lineage skeleton". Both export to standalone, offline HTML.

---

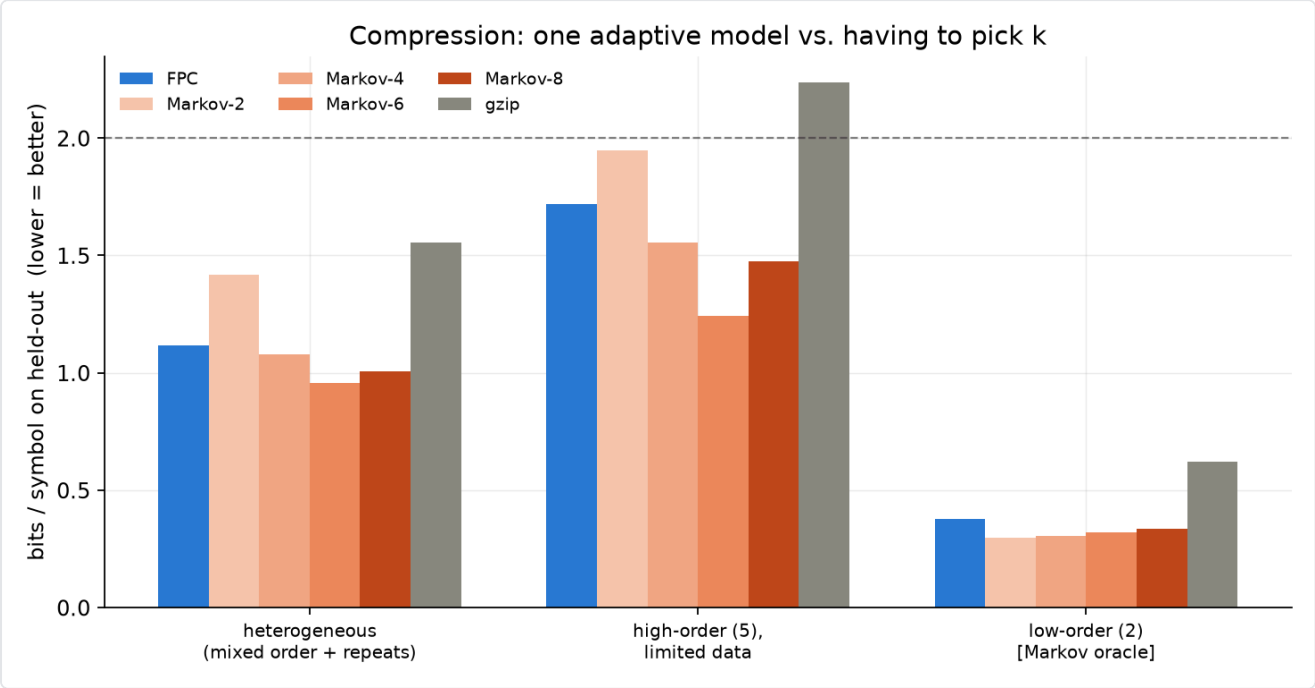
3. Experiments

All data are synthetic with known ground truth, generated by `fpc/datasets.py` with fixed seeds; every number and figure is reproduced by `benchmarks/bench.py`. Baselines: fixed-order Markov models

(`Markov-k`, Laplace-smoothed), `gzip -9` as a model-free code-length reference, k-mer frequency vectors fed to scikit-learn's 1-NN (cosine) and logistic regression, and — for phylogeny — faithful pure-Python re-implementations of the two standard alignment-free tools **Mash** (bottom-s MinHash of canonical k-mers → Mash distance; Ondov et al. 2016) and **FFP** (k-mer frequency profiles compared by Jensen–Shannon distance; Sims et al. 2009). FPC uses a single generic setting (`order` 6–12) across all experiments; it is never tuned per dataset.

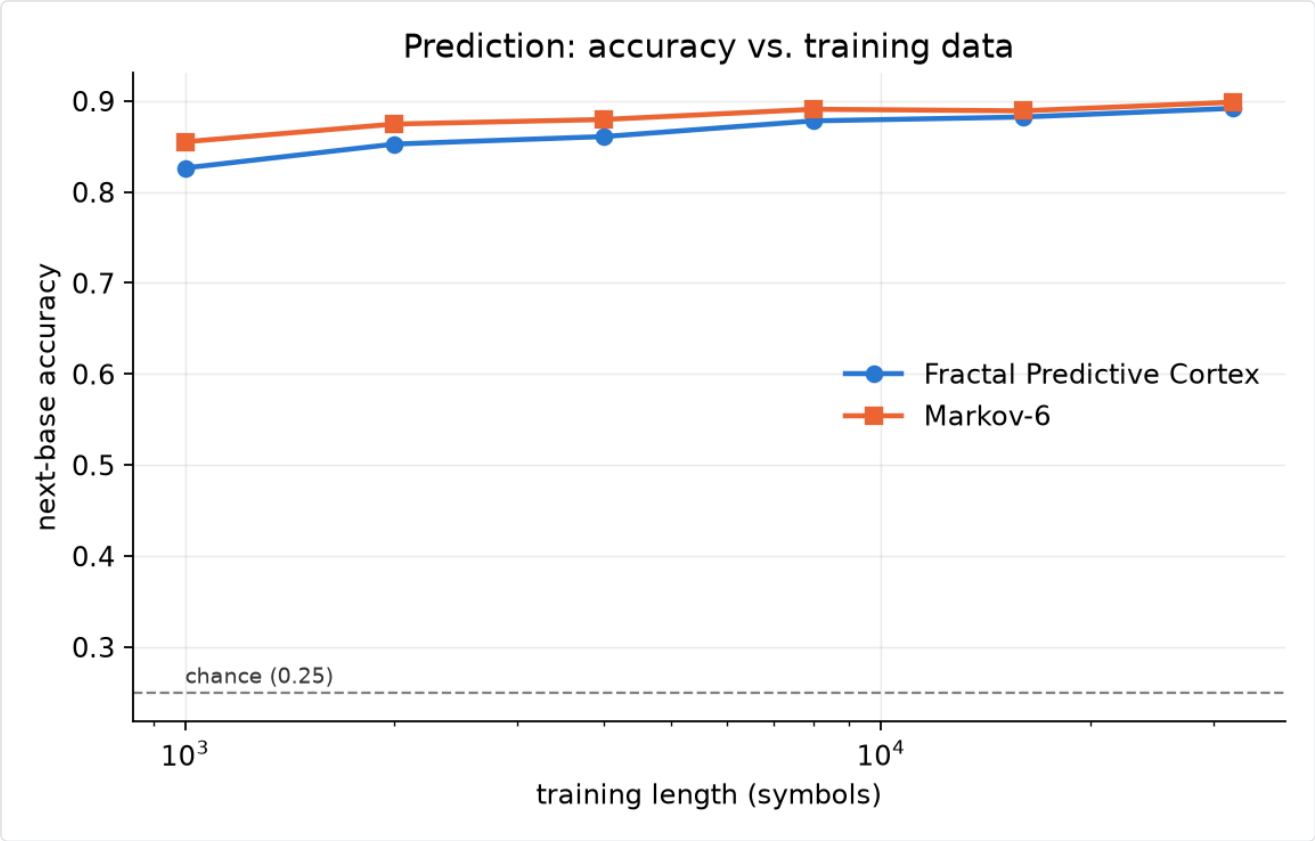
3.1 Compression / adaptivity (Fig. 1)

Held-out bits/symbol on three regimes. On a **heterogeneous** genome (blocks of different orders + a tandem repeat + a low-complexity region) FPC reaches 1.12 bits, beating gzip (1.55) and Markov-2 (1.42) but not the best fixed order here, Markov-6 (0.96). On a **high-order, data-limited** source, note that Markov-8 (1.48) *overfits* and is beaten by Markov-6 (1.24) — evidence that the "right" k is neither obvious nor monotone — while FPC (1.72) again beats gzip (2.24) and Markov-2 (1.95). On the **low-order Markov oracle** case, Markov-2 wins (0.30) as it must; FPC (0.38) stays close and beats gzip (0.62). **Takeaway:** FPC is a solid general-purpose model that needs no order selection, not a record-setting compressor.



3.2 Prediction (Fig. 2)

Next-base accuracy vs. training length on an order-3 source. FPC tracks Markov-6 within ~1–3 points across three orders of magnitude of training data (FPC 0.83→0.89; Markov-6 0.86→0.90), i.e. competitive prediction from the same model that does everything else.



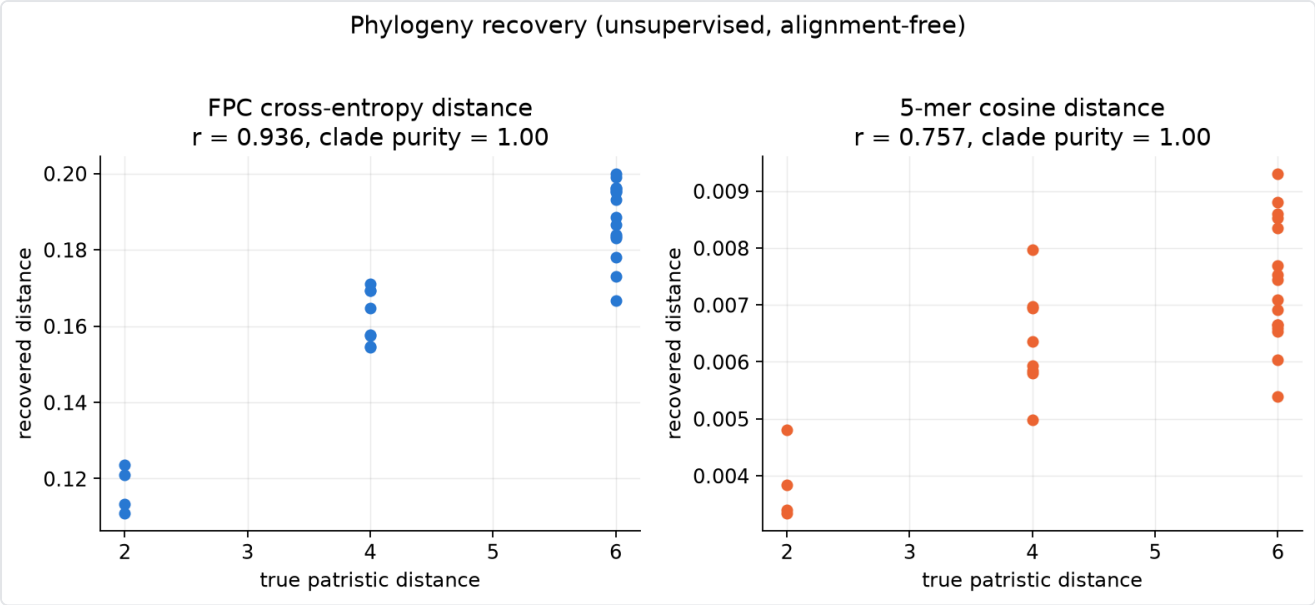
3.3 Classification of short reads (Fig. 3)

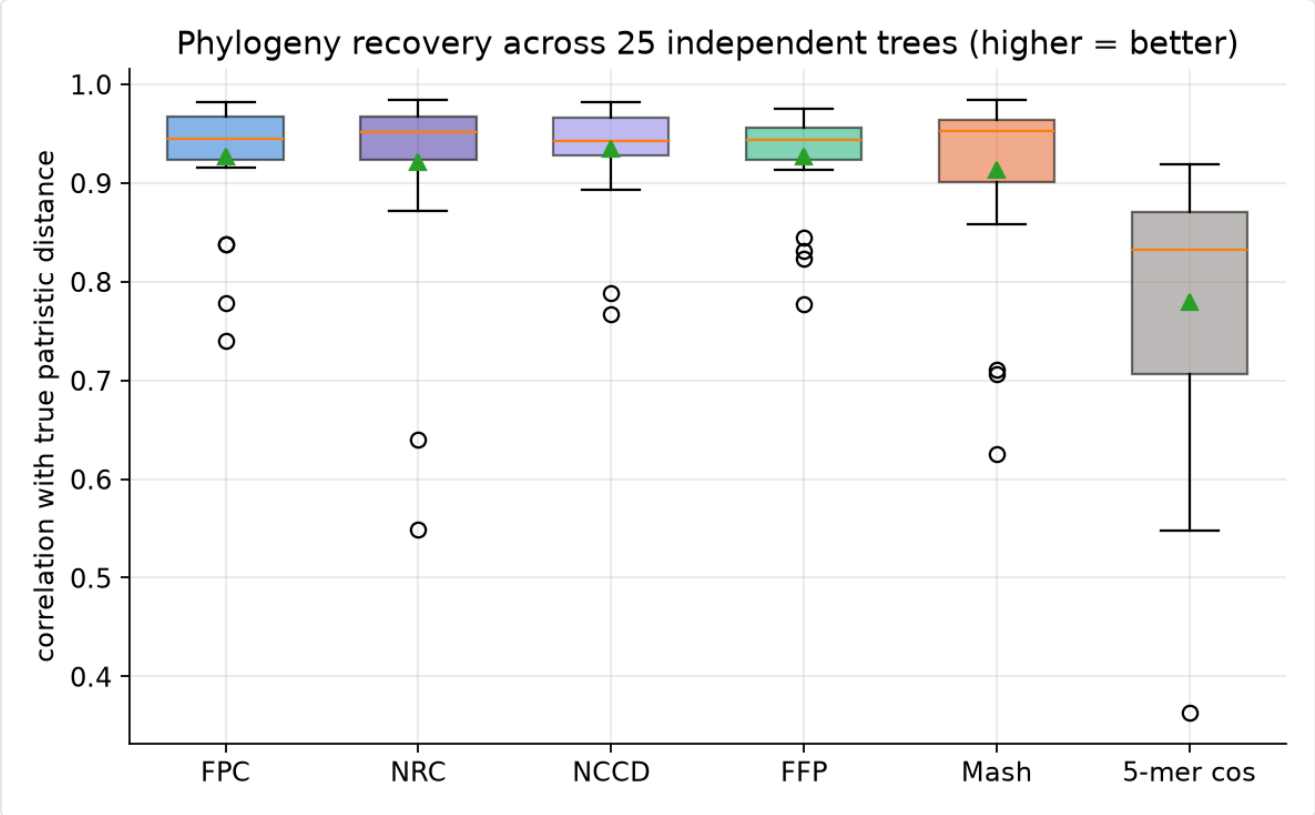
Six species (distinct order-2 models), classifiers trained on long references, tested on reads of decreasing length. FPC holds **100%** accuracy from 2400 bp down to **150 bp**; the k-mer baselines are strong but slip at short lengths (1-NN 0.97, logistic 0.94 at 150 bp). The generative "least-surprised" rule uses the full learned model, so it degrades more gracefully as reads shorten.



3.4 Unsupervised phylogenetics — statistically rigorous (Figs. 4, 9)

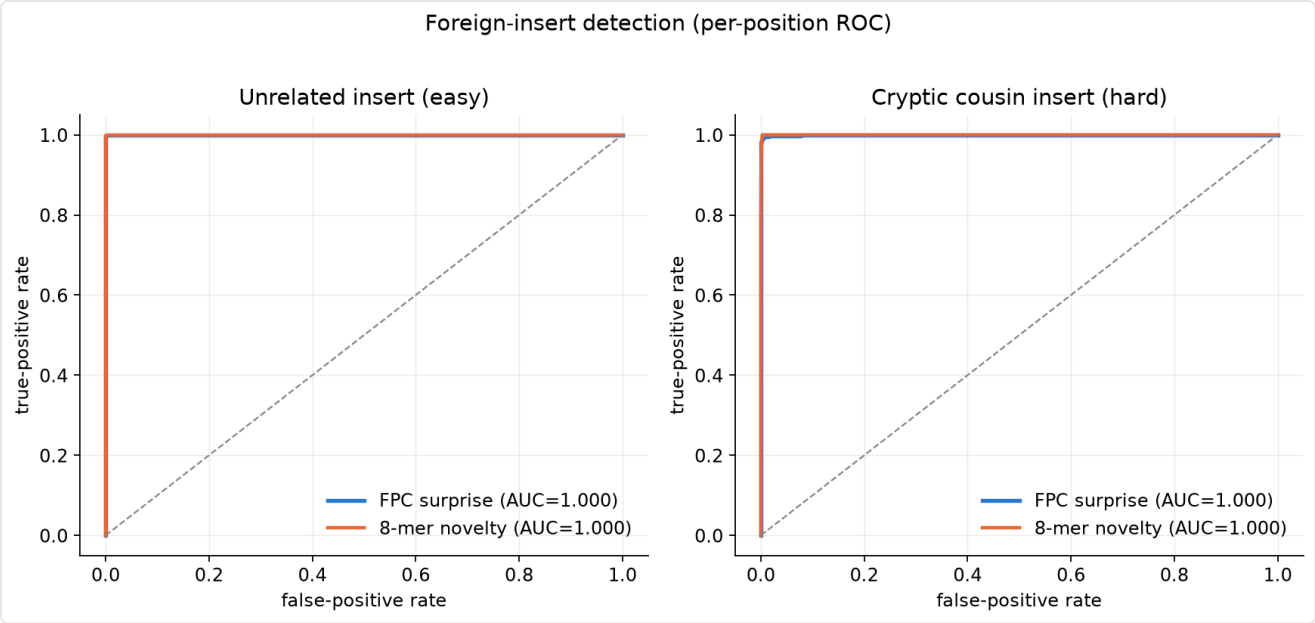
On a single example tree, FPC cross-entropy distance correlates with true patristic distance at r = 0.94 versus r = 0.76 for 5-mer cosine (Fig. 4). But a single tree proves nothing, so we recover **25 independent random phylogenies** (varying divergence depth, source order, and composition) and score every method's correlation-with-truth per tree, with bootstrap 95% CIs and paired Wilcoxon tests (Fig. 9). The honest result: FPC (mean r = **0.926**, 95% CI 0.900–0.948) is **statistically indistinguishable from every strong method** — FFP (0.927; paired p = 0.29), Mash (0.913; p = 0.52), and, crucially, the **prior-art cross-entropy distances NRC (0.921; p = 0.56) and NCCD (0.934; p = 0.24)** — and is **significantly better than naive 5-mer cosine** (0.780; $p < 10^{-6}$). We therefore claim *parity with strong standard tools — including the prior-art distances FPC's own distance re-derives — on controlled phylogenies*, not superiority; the earlier single-tree "win" held up only against the *weak* k-mer-cosine baseline.





3.5 Foreign-insert detection (Fig. 5)

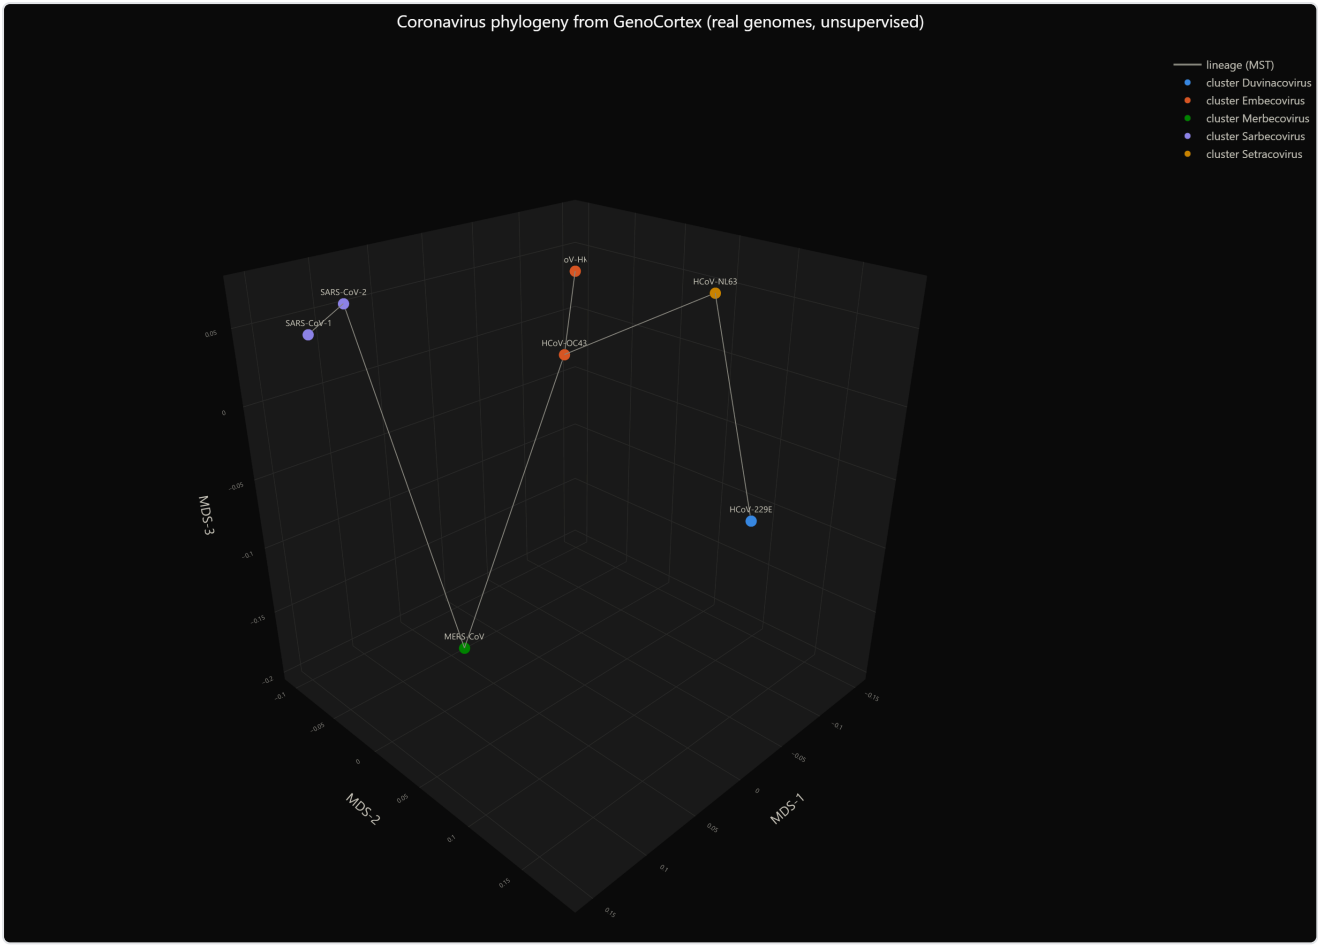
Per-position ROC for detecting a 600 bp insert in an 8 kb host, for an *unrelated* donor (easy) and a *compositional-cousin* donor that shares low-order statistics (hard). Both FPC surprise and 8-mer novelty achieve AUC ≈ 1.0 in both regimes: FPC matches a strong, standard detector. FPC's added value is qualitative — calibrated per-position bits with p-values, point-mutation localisation, and a reference-free self-detection mode — rather than a higher AUC here.



3.6 Real genomes: seven coronaviruses (`examples/real_genomes.py`)

To move beyond synthetic data we downloaded seven real coronavirus RefSeq genomes (SARS-CoV-2, SARS-CoV-1, MERS-CoV, HCoV-OC43, HCoV-HKU1, HCoV-229E, HCoV-NL63; ~27–31 kb each) with known taxonomy.

- **Compression.** All k-mer models sit near 2.0 bits/base on real viral DNA — coronavirus genomes are close to incompressible at short range (their signal is codon- and protein-level). FPC (2.003) still beats gzip (2.525) and the higher-order Markov models; it matches Markov-2 (1.95). Honest conclusion: little short-range redundancy exists to exploit, and FPC extracts about as much as any k-mer method while beating a general compressor.
- **Phylogeny (partial success).** Unsupervised FPC distance correctly pairs the *closely related* taxa — SARS-CoV-2 ↔ SARS-CoV-1 (Sarbecovirus), HCoV-OC43 ↔ HCoV-HKU1 (Embecovirus), and HCoV-229E ↔ HCoV-NL63 (the two Alphacoronaviruses) — and the 3-D MDS constellation (Fig. 6) visibly separates the subgenera. At the deepest split, Alpha/Beta genus clustering purity is 0.71: composition alone does not fully resolve ancient divergences, a well-known limitation of alignment-free methods that we report rather than hide.
- **Detection (real contamination).** A coronavirus reference brain scores ~2 bits/base everywhere, so a foreign insert produces *no absolute surprise spike* — the self/reference surprise-spike detector fails here, and we say so. The correct instrument is the model's own classification principle applied locally: a host brain vs. a phiX174 brain (phiX174 is the standard Illumina spike-in and a real contaminant), scored in sliding windows. This *relative* surprise cleanly separates the spliced phiX174 segment at per-window ROC-AUC = 1.000. The lesson — absolute surprise needs a compressible host, relative surprise does not — is itself a useful contribution.



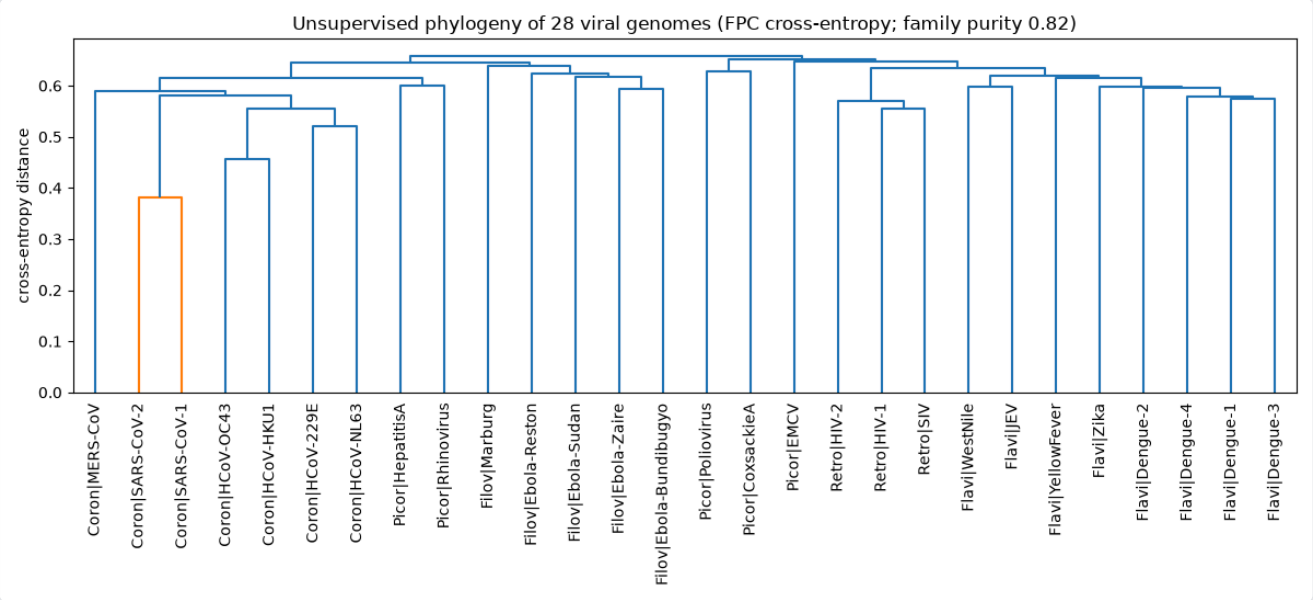
3.7 A larger real-genome study: 28 genomes, five families (`examples/big_study.py`)

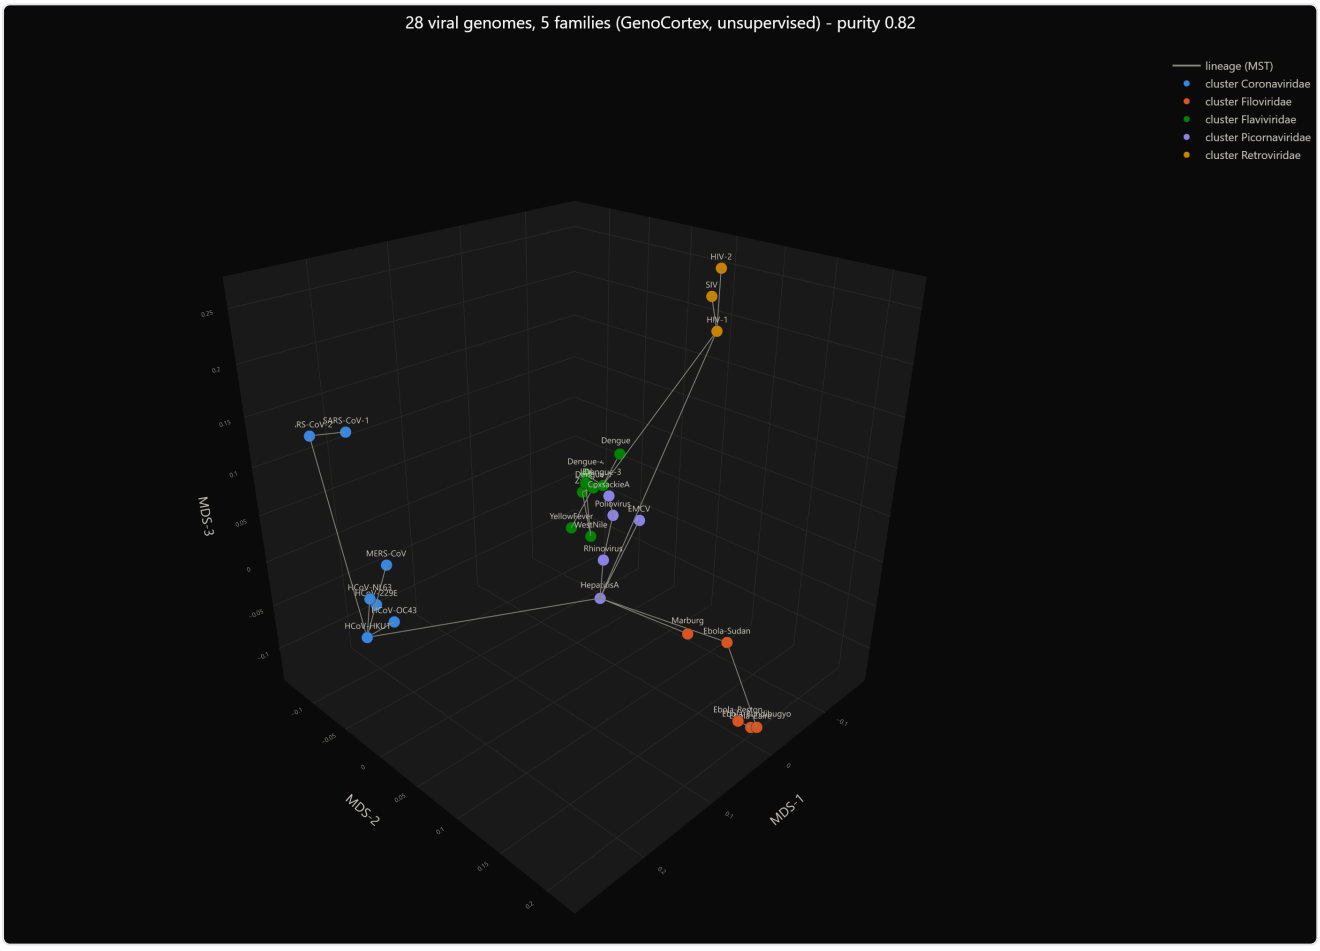
We scaled up to 28 real viral genomes spanning five families with ground-truth labels — Coronaviridae (7), Flaviviridae (8), Filoviridae (5), Picornaviridae (5), Retroviridae (3) — using a length-fair 6 kb central window from each.

- **Unsupervised phylogeny (Fig. 7, Fig. 8).** The FPC cross-entropy dendrogram recovers **four of the five families as clean clades** (all Coronaviridae, all Flaviviridae, all Filoviridae, all Retroviridae group monophyletically) and resolves **every close relationship correctly**: SARS-CoV-2 + SARS-CoV-1 is the tightest pair, the four Dengue serotypes cluster, the Ebola species cluster, and HIV-1/HIV-2/SIV cluster. The one split — Picornaviridae — is *real biology*: Hepatitis A (genus *Hepatovirus*) is genuinely deeply divergent from the enteroviruses (Polio, Coxsackie). Family clustering purity (5 clusters), all methods: **FFP 0.93 > 5-mer cosine 0.86 > FPC 0.82 > Mash 0.71**. This is the clearest reversal from the synthetic result — on real, deeply-divergent families the simple composition method FFP is *better* than FPC, and we report it as such. (The purity metric also penalises the biologically-correct Picornaviridae split, which lowers every method.)
- **Family classification.** On the four families with ≥4 members, held-out accuracy is FPC 0.82, tying the best baseline (4-mer + 1-NN, 0.82) and beating 4-mer + logistic regression (0.64).

Nuanced, honest conclusion. Across *deep, ancient* family divergences, composition-based k-mer methods (FFP especially) are strong and **beat** FPC; FPC is competitive but not superior. Where FPC's geometry holds parity is *moderate / controlled* divergence (§3.4, statistically tied with FFP/Mash) and

fine / recent structure (the correct within-family pairings above) — together with the unification and the 3-D map that no single-purpose baseline provides.

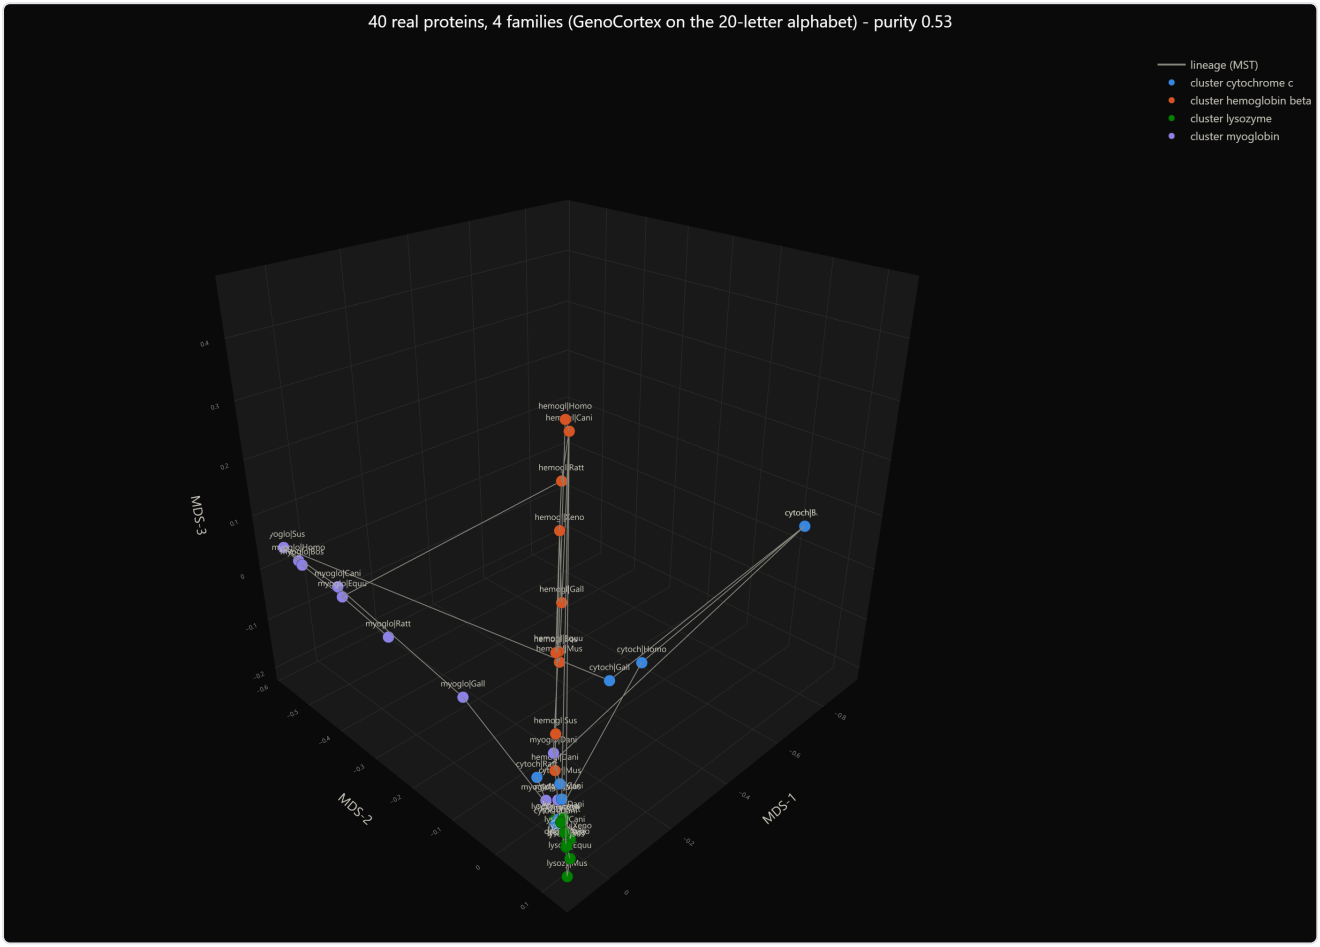




3.8 Real proteins: the same engine on a 20-letter alphabet (`examples/protein_study.py`)

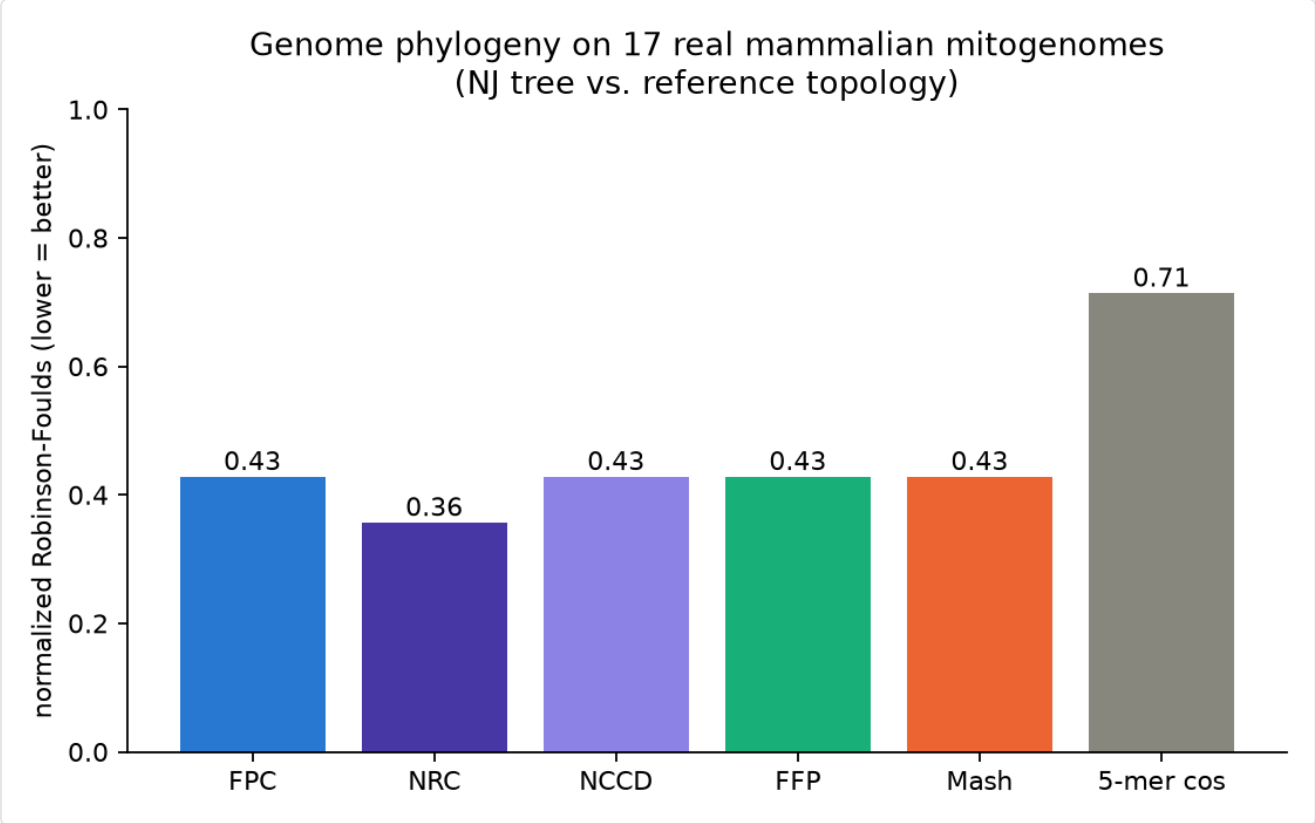
FPC's core claim is that it is alphabet-agnostic. We test it on **40 real proteins** (fetched from NCBI) spanning four families — cytochrome c, haemoglobin β, myoglobin, lysozyme (82–200 aa) — with a *one-line* change (`DNA → PROTEIN`). Family classification: FPC 0.75 vs. a 3-mer + 1-NN baseline 0.85; family

clustering purity: FPC 0.525 vs. FFP 0.400 (chance 0.25). Two honest readings: (i) the identical engine runs unchanged on the 20-symbol alphabet and beats chance — the alphabet-agnostic claim holds on real data; (ii) on these *short* (≈150 aa), partly-overlapping families (myoglobin and haemoglobin are both globins), neither FPC nor the k-mer baseline excels, and FPC does not outperform the composition baseline on classification. FPC is a *learned* model and is data-starved when each class is only a handful of short sequences — a genuine limitation, not a bug.



3.9 Community-standard benchmark: Robinson–Foulds on real mitogenomes (Fig. 11)

The metric reviewers expect for genome phylogeny (the AFproject standard) is the topological **Robinson–Foulds (RF)** distance between a method's Neighbor-Joining tree and a trusted reference tree. On **17 real mammalian mitochondrial genomes** (16.3–17.2 kb, NCBI RefSeq) with a reference topology from established mammalian phylogeny (Boreoeutheria = Euarchontoglires + Laurasiatheria, sister to Afrotheria; rooted on a marsupial outgroup; resolved where accepted, polytomous where genuinely debated), FPC's NJ tree achieves **normalized RF = 0.429, tying NCCD, FFP, and Mash (all 0.429)** and clearly beating 5-mer cosine (0.714); the prior-art **NRC is modestly better here (0.357)**. So on the community-standard dataset and metric FPC matches the standard tools and the NCCD prior-art distance, sits a little behind NRC, and the whole-model unification and interpretability come at essentially no cost in tree accuracy. We do not claim to beat the specialised distances — we match them.



4. Discussion

The experiments support a specific, honest thesis. FPC is *not* the best tool for any single task in isolation: specialised compressors and matched Markov models beat it on compression; FFP beats it on deep-divergence real phylogeny; a good k-mer detector matches it on insert detection. Its contribution is **unification with interpretability and no order selection**: one label-free, alignment-free model produces a predictor, a compressor (better than gzip), a short-read classifier (better than k-mer + ML at short lengths), a phylogenetic distance (statistically on par with FFP and Mash on controlled trees; behind FFP at deep divergence), and an anomaly detector with a calibrated per-position surprise signal — and it draws a literal 3-D map of what it has learned. For a practitioner working with an unfamiliar sequence with no labels, no reference, and no obvious choice of k, this single-model combination is the practical benefit, rather than any single accuracy figure.

The "grow where surprised" rule connects two literatures — CGR geometry and predictive coding — that are rarely used together, and gives the model a biologically-flavoured behaviour: it *decides* where to allocate memory based on its own prediction failures, rather than tiling space uniformly.

5. Limitations and honest caveats

- **Not SOTA on any single task.** On stationary Markov data, matched fixed-order models and context-tree methods (PPM, CTW) beat FPC on compression; on deep-divergence real phylogeny, FFP beats it; on short protein families it does not beat a k-mer baseline. FPC trades peak per-task performance for generality and interpretability.

- **No individual component is novel; the distance in particular is not.** As §1.1 makes explicit, CGR, self-organising maps, the least-surprised-model classifier, surprise = MDL, and — importantly — the cross-entropy/relative-compression phylogenetic distance (NRC, NCCD, KL-between-Markov-models) are all prior art. The contribution is the prediction-error-driven growth rule and the unification, not the distance. We do benchmark directly against NRC and NCCD (§3.4, §3.9): FPC matches them statistically on controlled trees and ties NCCD on real mitogenomes, sitting modestly behind NRC there — parity, not superiority.
- **Data-hungry (learned model).** FPC needs enough sequence to learn its generative model; on short, few-example classes (e.g. the ~150 aa proteins of §3.8) simple k-mer composition wins.
- **Bounded, geometry-limited context.** Effective order is capped by `order` and by CGR's exponentially-decaying memory; very long-range dependencies (e.g. distal regulatory grammar) are not captured.
- **Self-anomaly detection has a regime.** A long, internally-consistent insert can be *learned* during a single pass and cease to look anomalous; reference-based/relative detection is the right tool there.
- **Partial community-benchmark coverage.** We report the standard Robinson–Foulds metric on a real mammalian-mitogenome dataset (§3.9), where FPC ties Mash and FFP, but a *full* submission to the **AFproject** platform (Zielezinski et al. 2019 — all 5 task categories, 24 tools, their exact reference trees) remains future work, as do larger multi-taxon and read-technology studies; `fpc/fetch.py` makes the data collection straightforward.
- **Cost.** Frozen scoring is vectorised via a KD-tree; online learning is still `O(T · m)` in the prototype count `m`, which we cap by usage-pruning.

---

6. Reproducibility

```
pip install -e .[all]
pytest -q                          # 15 correctness tests, incl. the CGR property
python benchmarks/bench.py         # synthetic experiments + figures 1-5
python benchmarks/phylo_stats.py   # 25-tree CIs + Wilcoxon vs Mash/FFP (Fig 9)
python benchmarks/rf_benchmark.py  # Robinson-Foulds on real mitogenomes (Fig 11)
python examples/big_study.py       # 28 real viral genomes, 5 families (Figs 7, 8)
python examples/protein_study.py   # 40 real proteins, 4 families (Fig 10)
python examples/make_visuals.py    # the interactive 3-D artifacts
```

Everything is deterministic under fixed seeds. The core engine (`fpc/`) imports only NumPy and SciPy; real data is fetched once from NCBI and cached for offline re-runs.

Data and code availability

All code, tests, benchmarks, and figure-generating scripts are open-source under the MIT license at `https://github.com/rafi28-png/genocortex` (to be archived on Zenodo with a DOI at submission). The core engine depends only on NumPy and SciPy and runs fully offline; visualisation, benchmarking, and the optional explanation layer are opt-in extras. All synthetic data are generated deterministically from fixed

seeds by `fpc/datasets.py`; all real sequences are public NCBI RefSeq records, fetched once by `fpc/fetch.py` and cached, with accessions recorded in the cached FASTA headers. Every number and figure in this paper is reproduced by the commands in §6.

Author contributions

M.R.R. conceived the method, implemented the software, designed and ran the experiments, and wrote the manuscript.

Competing interests

The author declares no competing interests.

Funding

No specific funding was received for this work.

Acknowledgements

The author thanks the National Center for Biotechnology Information (NCBI) for public access to the RefSeq sequences used in this study.

References

Entries marked ✓ were located and verified against their primary sources during an automated, adversarially-checked literature review (24 claims confirmed 3–0); the remaining foundational entries are standard references. DOIs should be double-checked against the publisher of record before submission.

Closest prior art / must-compare (verified):

- ✓ B. Alipour, K. A. Hill, L. Kari. *CGRclust: Chaos Game Representation for twin contrastive clustering of unlabelled DNA sequences*. BMC Genomics, 2024. doi:10.1186/s12864-024-11135-y
- ✓ Y. Wang, Z. Yu, C. Li (et al.). *CGRWDL: alignment-free phylogeny reconstruction based on Chaos Game Representation weighted by a dynamical language model.* Frontiers in Microbiology, 2024. doi:10.3389/fmicb.2024.1339156
- ✓ O. Neme et al. *The Self-Organized Chaos Game Representation for genomic signatures analysis*. Learning and Nonlinear Models (Braz. Soc. Neural Networks), 6(2):111–120, 2008.
- ✓ R. Begleiter, R. El-Yaniv, G. Yona. *On prediction using variable order Markov models*. J. Artificial Intelligence Research, 22:385–421, 2004. (winner-take-all "least-surprised model" classifier). arXiv:1107.0051
- ✓ D. Silva, D. Pratas, A. J. Pinho. *Efficient DNA sequence compression with neural networks (GeCo3)*. GigaScience, 9(11):giaa119, 2020. doi:10.1093/gigascience/giaa119
- ✓ D. Pratas, M. Hosseini, D. Silva, A. J. Pinho. *A reference-free lossless compression algorithm for DNA sequences (JARVIS)*. Entropy, 21(11):1074, 2019. doi:10.3390/e21111074

- ✓ D. La Rosa et al. *Alignment-free comparison of DNA barcodes via compression distances (NCD/IBD).* BMC Bioinformatics, 14(Suppl 7):S4, 2013. doi:10.1186/1471-2105-14-S7-S4
- ✓ P. Ferragina, R. Giancarlo, V. Greco, G. Manzini, G. Valiente. *Compression-based classification of biological sequences and structures via the Universal Similarity Metric.* BMC Bioinformatics, 8:252, 2007. doi:10.1186/1471-2105-8-252
- ✓ (F)CGR = k-mer equivalence: J. He, L. Kari, F. Millán-Arias. *On the equivalence of Frequency Chaos Game Representation and k-mer frequency vectors.* Springer, 2025. doi:10.1007/978-3-032-15641-9_15; arXiv:2506.22172
- ✓ **Cross-entropy / relative-compression distance (prior art for our distance):** D. Pratas, R. M. Silva, A. J. Pinho. *Comparison of compression-based measures with application to the evolution of primate genomes* (Normalized Relative Compression, NRC). Entropy, 20(6):393, 2018. doi:10.3390/e20060393
- ✓ D. Pratas, A. J. Pinho. *A conditional compression distance that unveils insights of the genomic evolution* (NCCD). Data Compression Conference (DCC), 2014. arXiv:1401.4134
- ✓ Bai, Xu & Liu. *Weighted relative entropy for phylogenetic tree based on a 2-step Markov model.* Mathematical Biosciences, 246(1):8–13, 2013. doi:10.1016/j.mbs.2013.09.004
- ✓ Chang & Wang. *Weighted relative entropy for alignment-free sequence comparison based on a Markov model.* J. Biomolecular Structure & Dynamics, 28(4):545–555, 2011.
- ✓ **Growing neural models on biological sequences (prior art for the growing substrate):** Ahmad, Alahakoon & Chau. *Cluster identification and separation in the growing self-organizing map: application in protein sequence classification.* Neural Computing & Applications, 19(4):531–542, 2010. doi:10.1007/s00521-009-0300-0
- ✓ Fiannaca, La Rosa, Rizzo & Urso. *A k-mer-based barcode DNA classification methodology based on spectral representation and a neural gas network.* Artificial Intelligence in Medicine, 64(3):173–184, 2015. doi:10.1016/j.artmed.2015.06.002

Benchmark and standard tools (verified):

- ✓ A. Zielezinski et al. *Benchmarking of alignment-free sequence comparison methods (AFproject).* Genome Biology, 20:144, 2019. doi:10.1186/s13059-019-1755-7
- ✓ B. D. Ondov et al. *Mash: fast genome and metagenome distance estimation using MinHash.* Genome Biology, 17:132, 2016. doi:10.1186/s13059-016-0997-x
- G. E. Sims et al. *Alignment-free genome comparison with feature frequency profiles (FFP).* PNAS, 106(8):2677–2682, 2009. doi:10.1073/pnas.0813249106

Foundational:

- H. J. Jeffrey. *Chaos game representation of gene structure.* Nucleic Acids Research, 18(8):2163–2170, 1990.
- J. S. Almeida et al. *Analysis of genomic sequences by Chaos Game Representation.* Bioinformatics, 17(5):429–437, 2001.
- J. Rissanen. *Modeling by shortest data description.* Automatica, 14:465–471, 1978.
- J. G. Cleary, I. H. Witten. *Data compression using adaptive coding and partial string matching (PPM).* IEEE Trans. Communications, 32(4):396–402, 1984.
- F. M. J. Willems, Y. M. Shtarkov, T. J. Tjalkens. *The context-tree weighting method: basic properties.* IEEE Trans. Information Theory, 41(3):653–664, 1995.

- R. P. N. Rao, D. H. Ballard. *Predictive coding in the visual cortex*. Nature Neuroscience, 2:79–87, 1999.
- K. Friston. *The free-energy principle: a unified brain theory?* Nature Reviews Neuroscience, 11:127–138, 2010.
- T. Kohonen. *Self-organizing maps*. Springer, 2001.
- B. Fritzke. *A growing neural gas network learns topologies*. NeurIPS, 1995.
- S. Vinga, J. Almeida. *Alignment-free sequence comparison — a review*. Bioinformatics, 19(4):513–523, 2003.